

PH125.9x Data Science Capstone - Kaggle Kickstarter Dataset Analysis and Predictions

Marina Ganopolsky

12/3/2020

Contents

Overview	3
Dataset	3
Methods and Analysis	7
Data Insights, Cleaning & Visualizations	7
Initial Insights	7
Pledge Information	8
Data Manipulation	9
Categories & Funds Pledged	12
Project States	15
Pledged vs Goal	15
Feature Correlation & Pairs	17
Functions of Note	17
Evaluation Models	18
1. Classification Trees	19
2. Random Forest	21
3. Naive Bayes	21
4. Generalized Linear Models (GLM)	22
4.11 Predictor - Category	22
4.2 2 Predictors - Category and Time In Days	22
4.3 3 Predictors - Category, Time In Days, and Country	23

4.4 4 Predictors - Category, Time In Days, Country and Goal Amount in \$USD .	23
4.5 5 Predictors - Category, Time In Days, Country, Goal Amount in \$USD, and Launched Month	24
4.6 6 Predictors - Category, Time In Days, Country, Goal Amount in \$USD, Launched Month, and Launched Day of the Week	24
4.7 7 Predictors - Category, Time In Days, Country, Goal Amount in \$USD, Launched Month, Launched Day of the Week, and Currency	25
4.8 8 Predictors - Category, Time In Days, Country, Goal Amount in \$USD, Launched Month, Launched Day of the Week, Currency, and Year Launched.	25
Results	26
Conclusion	27
Report Summary	27
Other Options	27
Exclusion	27
Possible Extensions and Improvements	27
Appendix - Environment	28

Overview

This is Marina Ganopolsky's implementation of the **Choose Your Own Project** of the **Harvard: PH125.9x Data Science, Summer 2020: Capstone** course.

The objectives of this project are:

1. To find a dataset on which to conduct Exploratory Data Analysis (EDA), Feature Engineering, and on which to perform **machine learning tasks**.
2. To use as least 2 methods beyond linear regression during the project.
3. To achieve an accuracy of **> 50%**, and validate this with the **test data set**, as described below.
4. To summarize the findings, draw conclusions, and provide further recommendations, if needed.

The dataset I have chosen is the **[Kickstarter Dataset from Kaggle]**(<https://www.kaggle.com/kemical/kickstarter-projects?select=ks-projects-201801.csv>). This dataset represents Kickstarter projects between the years of **2009 and 2018**, as well as their respective success and failure information. Information about the country of origin of the project, the category, the goal amount, currency, etc is also present. A quick summary of the data wrangling modifications will be provided, as well as some visual representations of patterns present in the data, in the **Method and Analysis** section. This will inform the reader about the various trends present in the dataset.

An exploratory analysis will be performed to generate the predicted successes levels of the Kickstarter dataset, with various models, to be specified; a best-performing algorithm will be chosen, and the results of the calculations will be analyzed and explained; finally a conclusion will be provided.

Several models will be developed, ranging from the most naive to the most complex. These will be tested and compared using their **accuracy** value in order to assess their quality. In contrast to the **MovieLens** project there is no ideal evaluation criteria for this project. However, since we have several models available, and we will pick the model with 1. Over 50% accuracy (since that is better than simply random) 2. The highest accuracy as the chosen model to produce the accuracy of the model.

The data is broken into a **training dataset** (80%), and a **testing dataset**(20%). After evaluating every available algorithm I have come up with, the best-resulting model will be used on the **test dataset** to evaluate the quality of the predictions of the kickstarter project success.

This project can be found on **GitHub** here .

Dataset

The **Kickstarter** dataset is automatically downloaded with the code provided; The data sets can be additionally downloaded here:

- <https://raw.githubusercontent.com/mganopolsky/kickstarter/master/data/ks-projects-201801.csv>

As specified by the project description (and the provided code) the data is split into a **test set (20%)** and the **training set(80%)**, after some feature engineering is applied to it. The calculations and algorithm verification are done on the training set, and each algorithm is tested on the test set.

As a first step, in order to get some basic information about the data we're working with, we examine the **kickstarter** dataset. The original dataset contains 15 variables :

```
glimpse(data)
```

```
Rows: 378,661
Columns: 15
$ ID                  <dbl> 1000002330, 1000003930, 1000004038, 1000007540, 10...
$ name                <chr> "The Songs of Adelaide & Abullah", "Greeting From ...
$ category             <chr> "Poetry", "Narrative Film", "Narrative Film", "Mus...
$ main_category        <chr> "Publishing", "Film & Video", "Film & Video", "Mus...
$ currency              <chr> "GBP", "USD", "USD", "USD", "USD", "USD", "USD", ...
$ deadline              <date> 2015-10-09, 2017-11-01, 2013-02-26, 2012-04-16, 2...
$ goal                 <dbl> 1000, 30000, 45000, 5000, 19500, 50000, 1000, 2500...
$ launched              <dttm> 2015-08-11 12:12:28, 2017-09-02 04:43:57, 2013-01...
$ pledged               <dbl> 0.00, 2421.00, 220.00, 1.00, 1283.00, 52375.00, 12...
$ state                 <chr> "failed", "failed", "failed", "failed", "canceled"...
$ backers                <dbl> 0, 15, 3, 1, 14, 224, 16, 40, 58, 43, 0, 100, 0, 0...
$ country                <chr> "GB", "US", "US", "US", "US", "US", "US", "US", "U...
$ `usd pledged`          <dbl> 0.00, 100.00, 220.00, 1.00, 1283.00, 52375.00, 12...
$ usd_pledged_real       <dbl> 0.00, 2421.00, 220.00, 1.00, 1283.00, 52375.00, 12...
$ usd_goal_real           <dbl> 1533.95, 30000.00, 45000.00, 5000.00, 19500.00, 50...
```

A fair amount of data manipulation has been done to the datasets created with the provided code. The changes include:

- **New Field** : Adding a time interval in days, as the difference between the launch date and the deadline date : **time_int**
- **New Field** : Extracting the month of the year from the launch date, as saving it as a factor in **launched_month**
- **New Field** : Extracting the year from the launch date, as saving it as a factor in **launched_year**
- **New Field** : Extracting the day of the week from the launch date, as saving it as a factor in **launched_day_of_week**

- **New Field** : Added a field as a rounded numeric representing the pledge:goal ratio in USD **pledged_ratio**
- **New Field** : Added a field as a rounded numeric representing the average pledge per backer in USD **avg_backer_pldg**
- Changed the **launched** field into a date (as opposed to a date/time field)
- **Removed Field** : Removed the **usd pledged** field, as there was the usd_pledged_real represented what was actually needed.
- After the data wrangling and data visualization exercises (see below) are complete, the dataset is split up into :
 - **A training dataset** - the data the algorithms will be trained on, **80%** of the remaining data.
 - **A testing dataset** - the data the algorithms will be tested on, **20%** of the remaining data.

Per row, this is a representation of a single Kickstarter campaign.

As a first glance, here's a snapshot of dataset:

ID	name	category	main_category
Min. :5.971e+03	Length:378661	Length:378661	Length:378661
1st Qu.:5.383e+08	Class :character	Class :character	Class :character
Median :1.075e+09	Mode :character	Mode :character	Mode :character
Mean :1.075e+09			
3rd Qu.:1.610e+09			
Max. :2.147e+09			
currency	deadline	goal	launched
USD :295365	Min. :2009-05-03	Min. : 0	Min. :1970-01-01
GBP : 34132	1st Qu.:2013-06-08	1st Qu.: 2000	1st Qu.:2013-05-07
EUR : 17405	Median :2015-01-14	Median : 5200	Median :2014-12-10
CAD : 14962	Mean :2014-11-01	Mean : 49081	Mean :2014-09-28
AUD : 7950	3rd Qu.:2016-04-28	3rd Qu.: 16000	3rd Qu.:2016-03-24
SEK : 1788	Max. :2018-03-03	Max. :1000000000	Max. :2018-01-02
(Other): 7059			
pledged	state	backers	country
Min. : 0	Length:378661	Min. : 0.0	Length:378661
1st Qu.: 30	Class :character	1st Qu.: 2.0	Class :character
Median : 620	Mode :character	Median : 12.0	Mode :character
Mean : 9683		Mean : 105.6	
3rd Qu.: 4076		3rd Qu.: 56.0	
Max. :20338986		Max. :219382.0	

usd_pledged_real	usd_goal_real	time_int	pledged_ratio
Min. : 0	Min. : 0	Min. : 1.00	Min. : 0.00
1st Qu.: 31	1st Qu.: 2000	1st Qu.: 30.00	1st Qu.: 0.00
Median : 624	Median : 5500	Median : 30.00	Median : 0.13
Mean : 9059	Mean : 45454	Mean : 34.48	Mean : 3.24
3rd Qu.: 4050	3rd Qu.: 15500	3rd Qu.: 37.00	3rd Qu.: 1.07
Max. : 20338986	Max. : 166361391	Max. : 16739.00	Max. : 104277.89

avg_backer_pldg	launched_month	launched_day_of_week	launched_year
Min. : 0.00	07 : 36367	1:61228	2015 : 77300
1st Qu.: 13.00	03 : 33946	2:77307	2014 : 67745
Median : 40.00	10 : 33490	3:67438	2016 : 57184
Mean : 64.58	11 : 32890	4:60596	2017 : 52200
3rd Qu.: 75.00	05 : 32888	5:58327	2013 : 44851
Max. : 10000.00	06 : 32623	6:32717	2012 : 41165
	(Other): 176457	7:21048	(Other): 38216

The first few lines of the dataset look like this:

ID	name
1 1000002330	The Songs of Adelaide & Abullah
2 1000003930	Greeting From Earth: ZGAC Arts Capsule For ET
3 1000004038	Where is Hank?
4 1000007540	ToshiCapital Rekordz Needs Help to Complete Album
5 1000011046	Community Film Project: The Art of Neighborhood Filmmaking
6 1000014025	Monarch Espresso Bar

category	main_category	currency	deadline	goal	launched	pledged
1 Poetry	Publishing	GBP	2015-10-09	1000	2015-08-11	0
2 Narrative Film	Film & Video	USD	2017-11-01	30000	2017-09-02	2421
3 Narrative Film	Film & Video	USD	2013-02-26	45000	2013-01-12	220
4 Music	Music	USD	2012-04-16	5000	2012-03-17	1
5 Film & Video	Film & Video	USD	2015-08-29	19500	2015-07-04	1283
6 Restaurants	Food	USD	2016-04-01	50000	2016-02-26	52375

state	backers	country	usd_pledged_real	usd_goal_real	time_int
1 failed	0	GB	0	1533.95	59
2 failed	15	US	2421	30000.00	60
3 failed	3	US	220	45000.00	45
4 failed	1	US	1	5000.00	30
5 canceled	14	US	1283	19500.00	56
6 successful	224	US	52375	50000.00	35

pledged_ratio	avg_backer_pldg	launched_month	launched_day_of_week
1 0.00	0	08	2
2 0.08	161	09	6
3 0.00	73	01	6

```

4          0.00          1          03          6
5          0.07         92          07          6
6          1.05        234          02          5
launched_year
1        2015
2        2017
3        2013
4        2012
5        2015
6        2016

```

Methods and Analysis

Data Insights, Cleaning & Visualizations

Initial Insights

Glancing over the information most things seem fine at a first review; one thing that jumps out is the minimum value of the **launched** field, which is 1/1/1970. This is probably an error in the data - let's examine it:

```
ds %>% distinct(launched) %>% arrange(launched) %>% head()
```

```
# A tibble: 6 x 1
  launched
  <date>
1 1970-01-01
2 2009-04-21
3 2009-04-23
4 2009-04-24
5 2009-04-25
6 2009-04-27
```

Viewing a list of the furthest ordered launch dates, we see that 1/1/1970 is a strange outlier. We dig further by analyzing the entries with this information:

```
ds %>% filter(launched=='1970-01-01') %>%
  dplyr::select(launched, name, category, deadline, backers, goal)
```

```
# A tibble: 7 x 6
  launched    name      category    deadline    backers    goal
  <date>     <chr>     <chr>       <date>      <dbl>     <dbl>
1 1970-01-01  " "        " "        " "        0.00      0.00
```

```

1 1970-01-01 "Salt of the Earth: A Dead Sea ~ Film & V~ 2010-09-15      0  5000
2 1970-01-01 "1st Super-Size Painting - Soci~ Art      2010-08-14      0 15000
3 1970-01-01 "\"ICHOR\" (Canceled)"      Film & V~ 2010-05-21      0   700
4 1970-01-01 "Support Solo Theater! Help \"U~ Theater 2010-06-01      0  4000
5 1970-01-01 "Help RIZ Make A Charity Album:~ Music 2010-05-04      0 10000
6 1970-01-01 "Identity Communications Infogr~ Design 2010-04-10      0   500
7 1970-01-01 "Student Auditions Music 2015"  Publishi~ 2015-10-31      0  1900

```

Based on the names, most of these campaigns have been **cancelled**. It's worth noting that in the predictive modeling section, our models will exclude projects that do not fall into the "**successful**" or "**failed**" categories. Therefore, we will remove these entries from the data.

Pledge Information

What kind of information can be gleaned from failed projects? We examine projects where the pledge ratio is less than 1 (meaning that they did not meet their fundraise goal.)

```

# A tibble: 6 x 1
  state
  <chr>
1 failed
2 canceled
3 live
4 undefined
5 suspended
6 successful

```

It appears that some projects are successful despite not having met their funding goal!

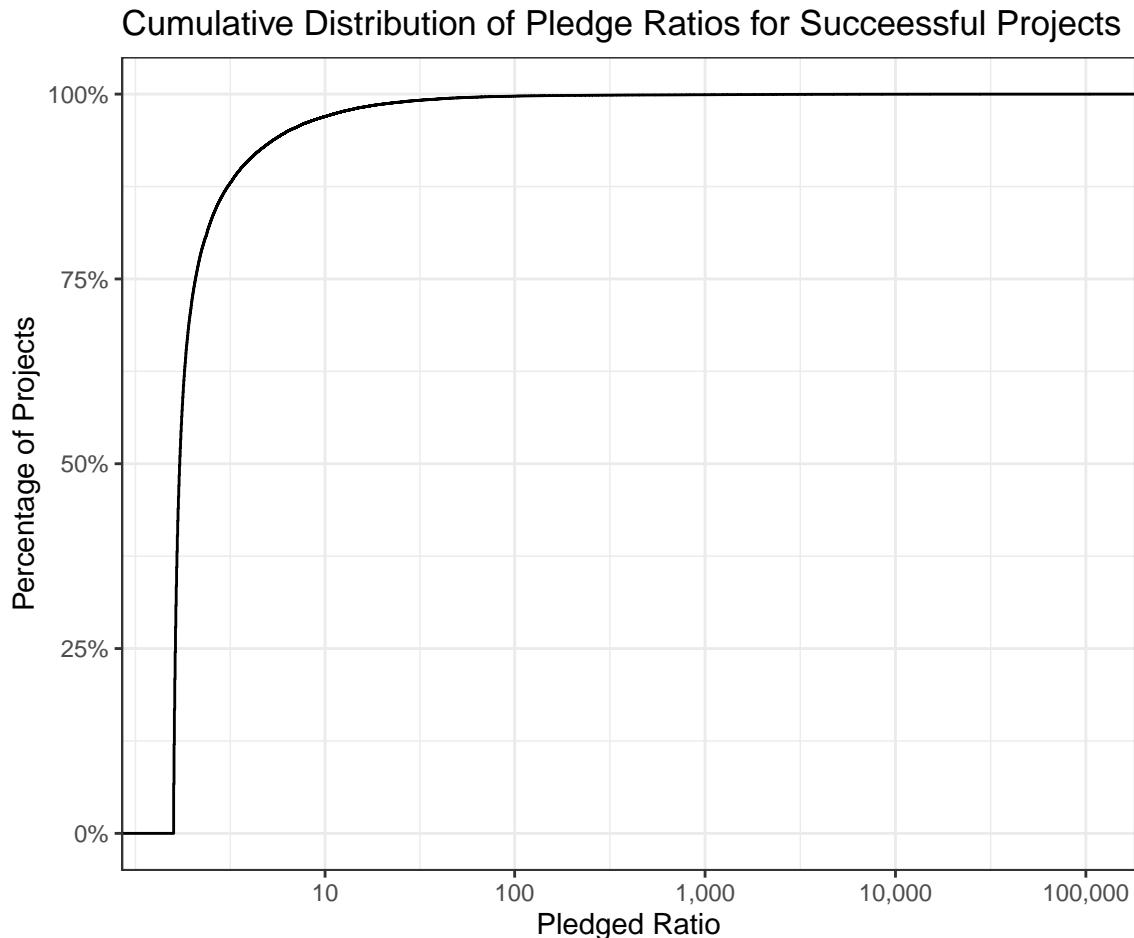
```

# A tibble: 1 x 20
  ID name  category main_category currency deadline      goal launched
  <dbl> <chr> <chr>    <chr>       <fct>     <date>      <dbl> <date>
1 1.77e9 Bord~ Shorts  Film & Video  USD      2015-12-27 36000 2015-11-12
# ... with 12 more variables: pledged <dbl>, state <chr>, backers <dbl>,
#   country <chr>, usd_pledged_real <dbl>, usd_goal_real <dbl>, time_int <dbl>,
#   pledged_ratio <dbl>, avg_backer_pldg <dbl>, launched_month <fct>,
#   launched_day_of_week <fct>, launched_year <fct>

```

The project raised a significant sum though - 85% of what it needed; and it went on to be successful. However, it's the only one of its kind. This of course, begs the question of what happens kind of funding rates we see for both failed and successful projects. We will examine this later in this document.

Viewing the pledge ratios' cumulative distribution, **it is obvious that most of the successful projects are funded around 100% or slightly higher - with the ratios hovering around 1. However, there are actually a significant amount of projects that are funded at thousands of times their goal value.**



Data Manipulation

Since we will be using models that require fields set as factors, the predictive fields of the dataset will be manipulated to be changed into factors.

The levels of the factors of these fields can be viewed here:

```
levels(ds$category)
```

```
[1] "3D Printing"          "Academic"           "Accessories"  
[4] "Action"              "Animals"             "Animation"  
[7] "Anthologies"         "Apparel"            "Apps"  
[10] "Architecture"        "Art"                "Art Books"
```

[13]	"Audio"	"Bacon"	"Blues"
[16]	"Calendars"	"Camera Equipment"	"Candles"
[19]	"Ceramics"	"Children's Books"	"Childrenswear"
[22]	"Chiptune"	"Civic Design"	"Classical Music"
[25]	"Comedy"	"Comic Books"	"Comics"
[28]	"Community Gardens"	"Conceptual Art"	"Cookbooks"
[31]	"Country & Folk"	"Couture"	"Crafts"
[34]	"Crochet"	"Dance"	"Design"
[37]	"Digital Art"	"DIY"	"DIY Electronics"
[40]	"Documentary"	"Drama"	"Drinks"
[43]	"Electronic Music"	"Embroidery"	"Events"
[46]	"Experimental"	"Fabrication Tools"	"Faith"
[49]	"Family"	"Fantasy"	"Farmer's Markets"
[52]	"Farms"	"Fashion"	"Festivals"
[55]	"Fiction"	"Film & Video"	"Fine Art"
[58]	"Flight"	"Food"	"Food Trucks"
[61]	"Footwear"	"Gadgets"	"Games"
[64]	"Gaming Hardware"	"Glass"	"Graphic Design"
[67]	"Graphic Novels"	"Hardware"	"Hip-Hop"
[70]	"Horror"	"Illustration"	"Immersive"
[73]	"Indie Rock"	"Installations"	"Interactive Design"
[76]	"Jazz"	"Jewelry"	"Journalism"
[79]	"Kids"	"Knitting"	"Latin"
[82]	"Letterpress"	"Literary Journals"	"Literary Spaces"
[85]	"Live Games"	"Makerspaces"	"Metal"
[88]	"Mixed Media"	"Mobile Games"	"Movie Theaters"
[91]	"Music"	"Music Videos"	"Musical"
[94]	"Narrative Film"	"Nature"	"Nonfiction"
[97]	"Painting"	"People"	"Performance Art"
[100]	"Performances"	"Periodicals"	"Pet Fashion"
[103]	"Photo"	"Photobooks"	"Photography"
[106]	"Places"	"Playing Cards"	"Plays"
[109]	"Poetry"	"Pop"	"Pottery"
[112]	"Print"	"Printing"	"Product Design"
[115]	"Public Art"	"Publishing"	"Punk"
[118]	"Puzzles"	"Quilts"	"R&B"
[121]	"Radio & Podcasts"	"Ready-to-wear"	"Residencies"
[124]	"Restaurants"	"Robots"	"Rock"
[127]	"Romance"	"Science Fiction"	"Sculpture"
[130]	"Shorts"	"Small Batch"	"Software"
[133]	"Sound"	"Space Exploration"	"Spaces"
[136]	"Stationery"	"Tabletop Games"	"Taxidermy"
[139]	"Technology"	"Television"	"Textiles"
[142]	"Theater"	"Thrillers"	"Translations"
[145]	"Typography"	"Vegan"	"Video"

```
[148] "Video Art"           "Video Games"        "Wearables"
[151] "Weaving"              "Web"                 "Webcomics"
[154] "Webseries"             "Woodworking"        "Workshops"
[157] "World Music"           "Young Adult"         "Zines"

levels(ds$main_category)

[1] "Art"                  "Comics"               "Crafts"                "Dance"                "Design"
[6] "Fashion"               "Film & Video"        "Food"                 "Games"                "Journalism"
[11] "Music"                 "Photography"          "Publishing"           "Technology"          "Theater"

levels(ds$currency)

[1] "AUD"      "CAD"      "CHF"      "DKK"      "EUR"      "GBP"      "HKD"      "JPY"      "MXN"      "NOK"      "NZD"      "SEK"
[13] "SGD"      "USD"

levels(ds$state)

[1] "canceled"   "failed"     "live"       "successful" "suspended"
[6] "undefined"

levels(ds$launched_month)

[1] "01" "02" "03" "04" "05" "06" "07" "08" "09" "10" "11" "12"

levels(ds$launched_day_of_week)

[1] "1" "2" "3" "4" "5" "6" "7"

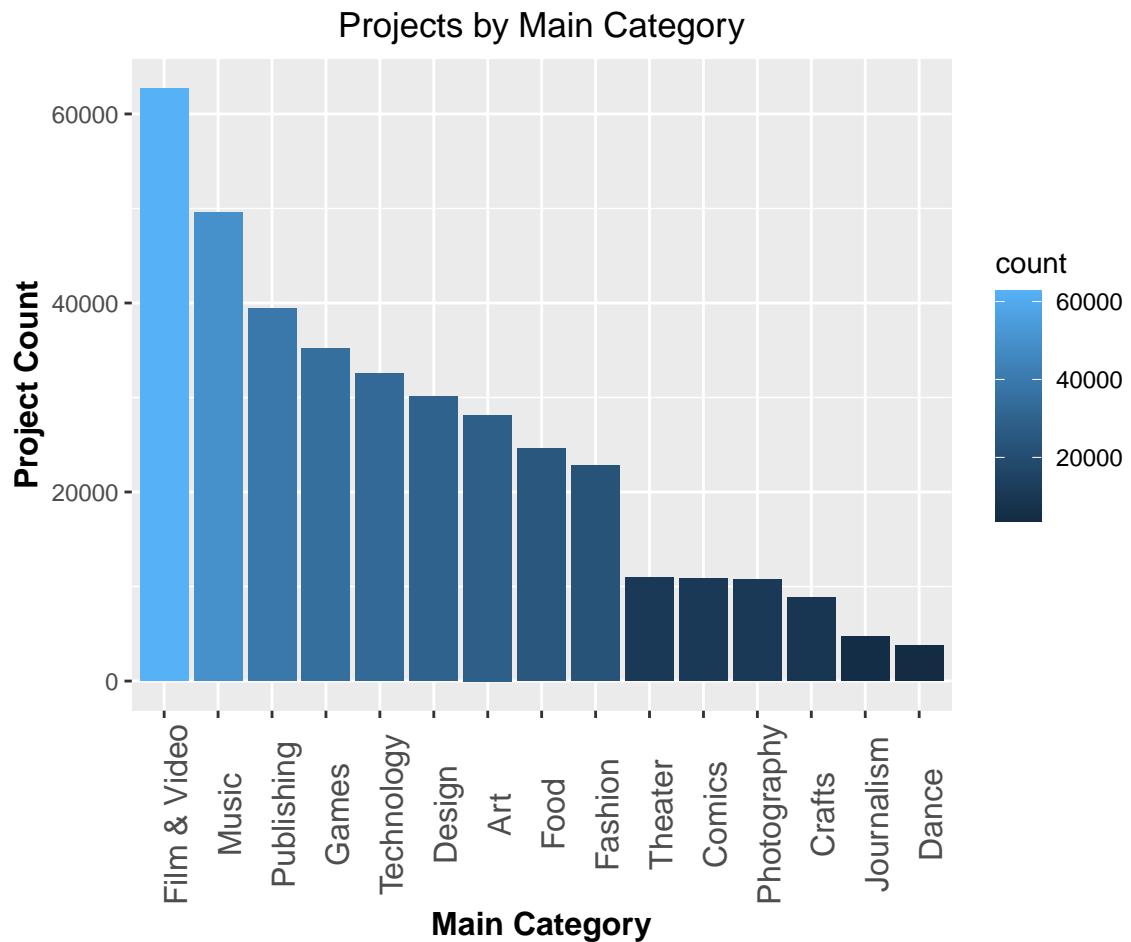
levels(ds$launched_year)

[1] "1970" "2009" "2010" "2011" "2012" "2013" "2014" "2015" "2016" "2017"
[11] "2018"
```

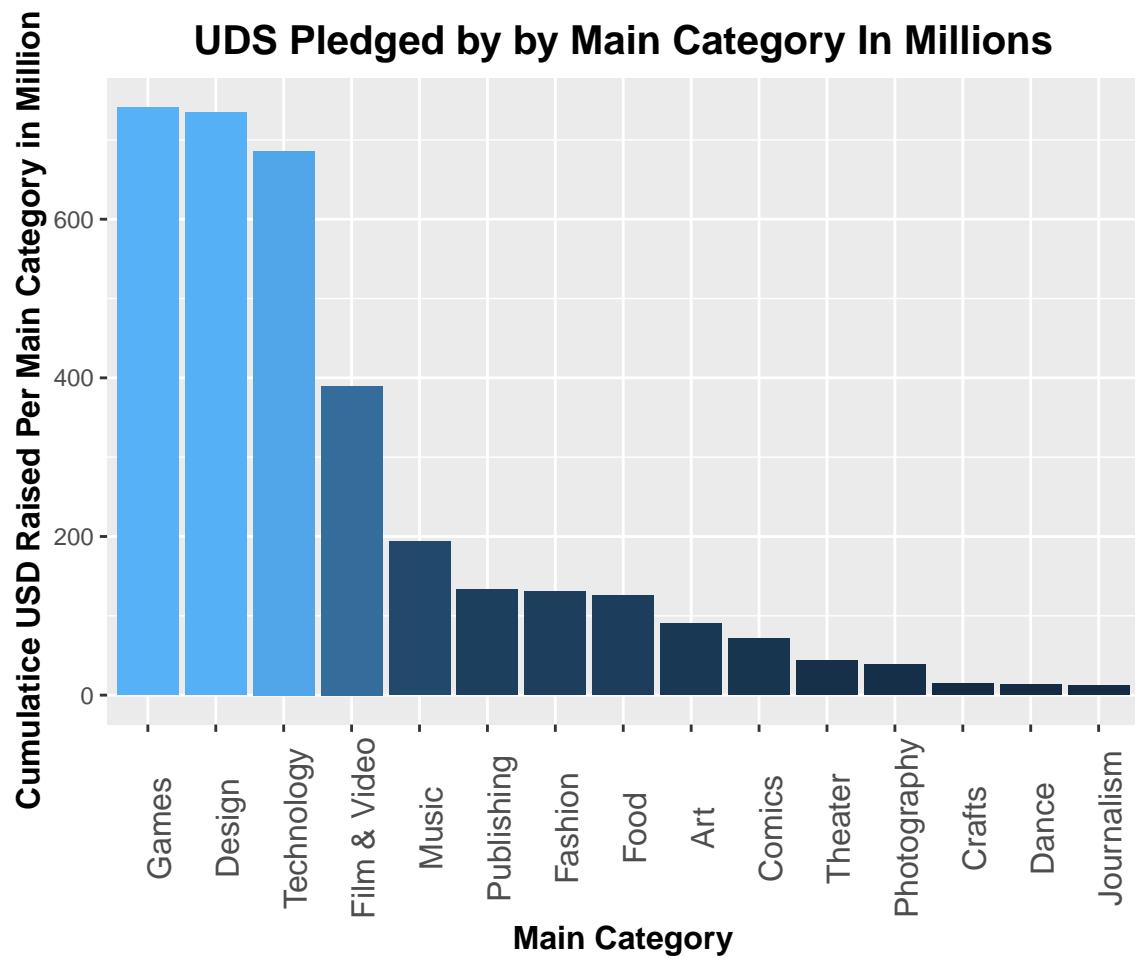
It now seems that the **country** field has some strange values (**"N,O"**) in it; we will remove these records.

Categories & Funds Pledged

Each project has a **main_category** and a **category** section. Which **main_category** do users attempt to fund most often?

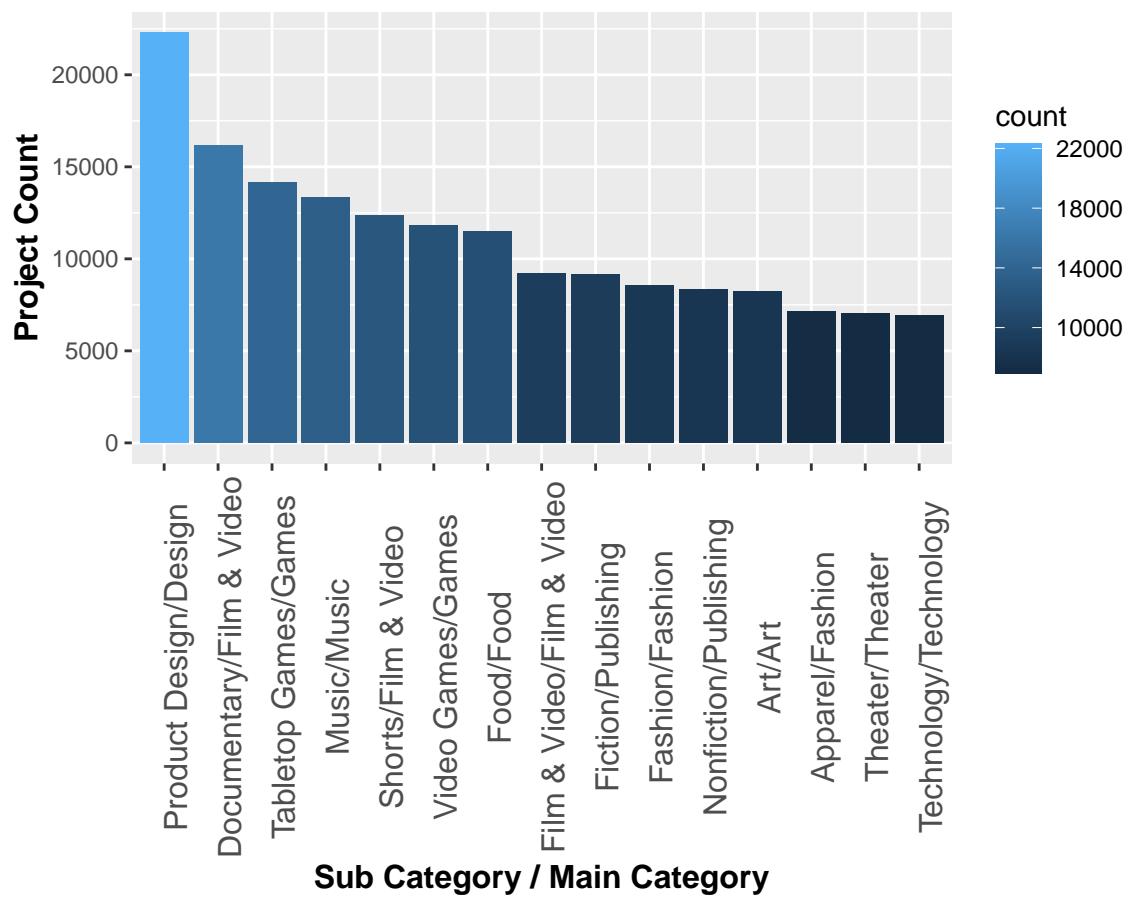


In terms of the main categories, **Film & video** projects seem to be the **most** widespread, while **Dance** is the **least**. What does this look like in terms of funds raised?

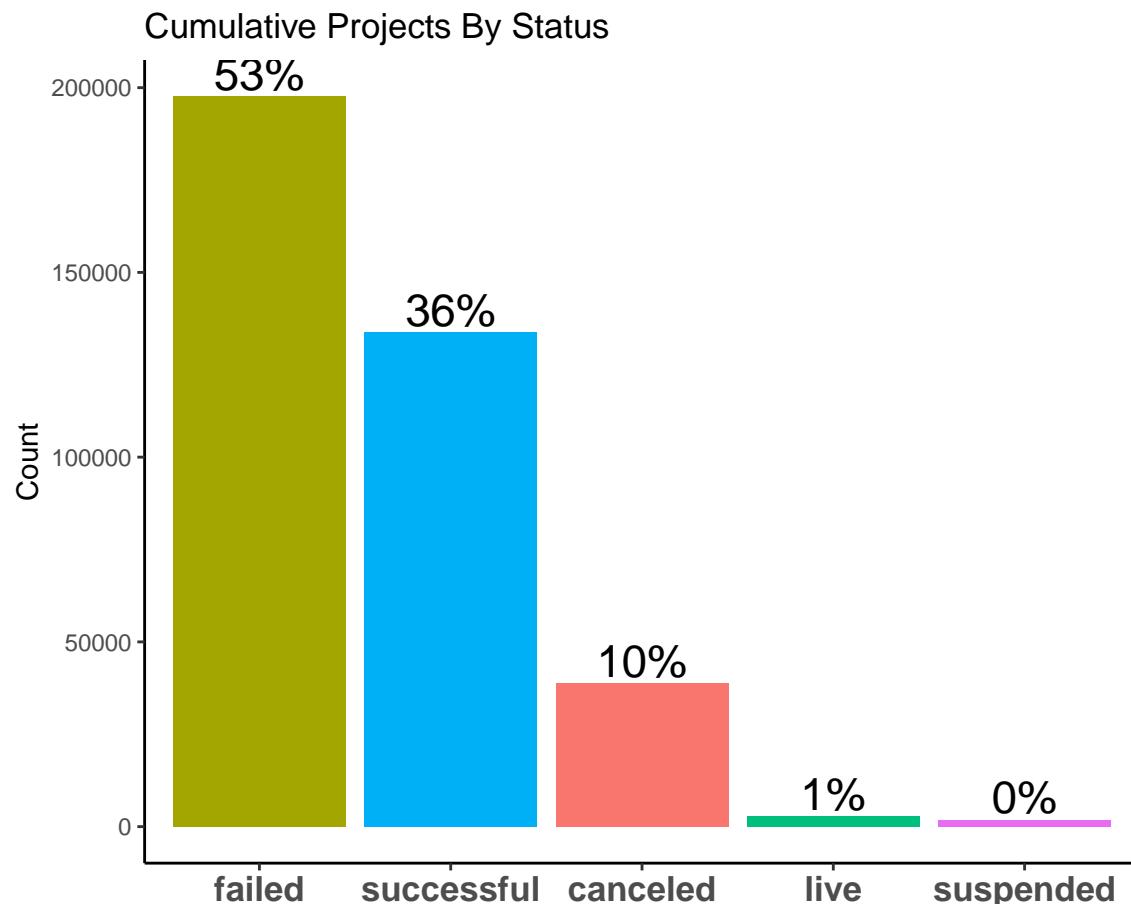


And what about sub-categories?

Projects by Sub / Main Category



Project States



From the data and the bar graph of the project states, we can make these conclusions:

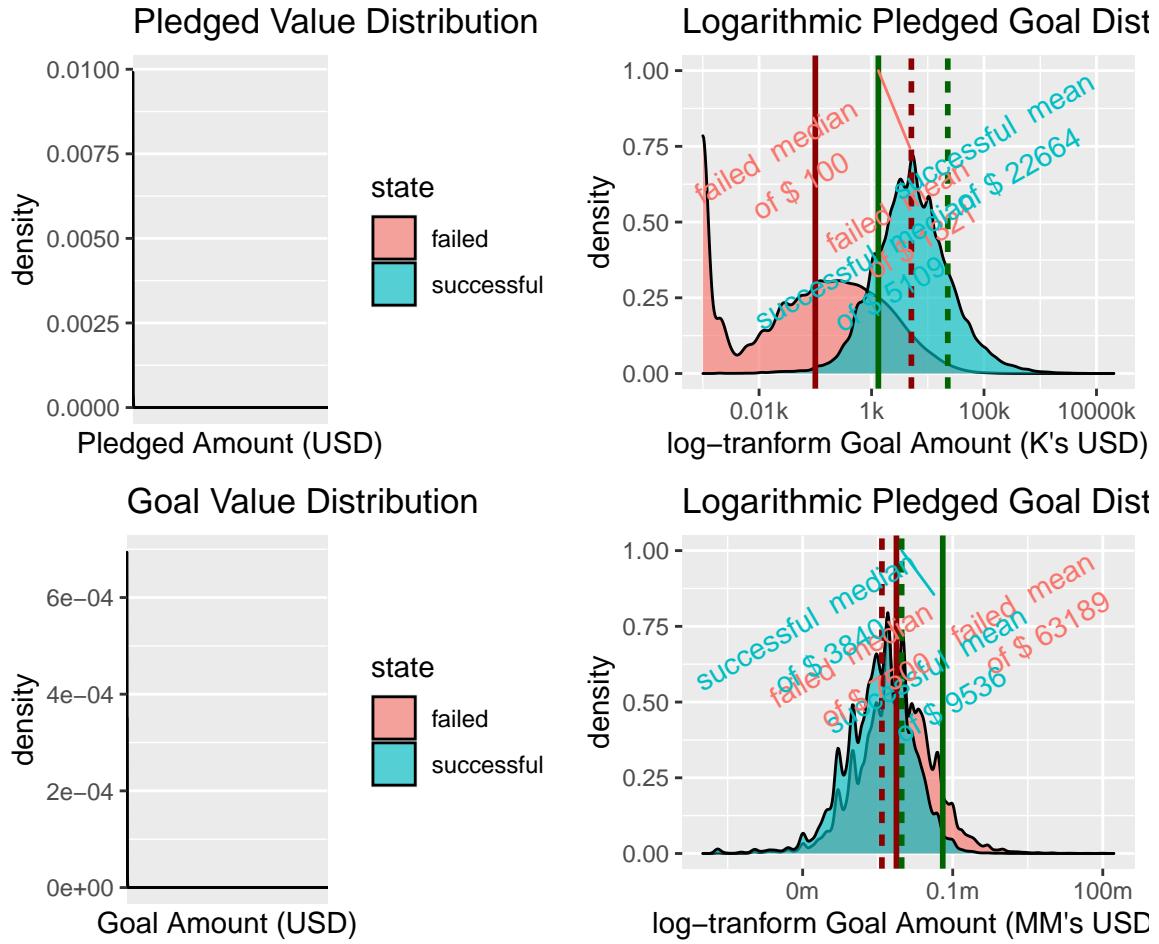
* The vast majority of the projects fall into the “**failed**” or “**successful**” category. * There are 6 discrete options. * Most projects **FAIL**.

The information in the dataset includes some obvious descriptors of what happens in a project, but is hardly a full picture of what happens during a fundraising campaign. We don’t know about the marketing efforts, time/funds spent on those, or anything else. Therefore, it’s hardly a detailed model we’re building, but it’s a start. As mentioned earlier, we will be focusing exclusively on the projects that fall into the “**failed**” or “**successful**” categories.

Pledged vs Goal

Let’s examine the distribution of **pledged value** and **goal value** in **successful** vs **failed projects**; As is evidenced by the density plots on the left side, little info can be gleaned here. However, once the visualization is **log-transformed**, the distributions of the successful projects seem to be nearing normal. The pledged values are normally distributed for successful projects, but not so much for the failed ones; this makes intuitive sense since many failed projects had raised no

money at all.



Digging into this further, a box plot of pledge ratios paints a very clear picture of failed vs. successful projects - and this makes sense intuitively, as well. We've already seen that most successful projects are funded at over 100\$ of goal.

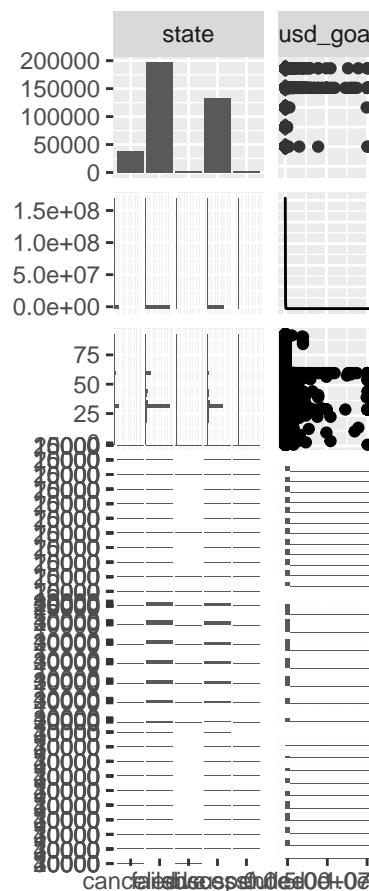
10%	20%	30%	40%	50%	60%	70%	80%	90%
0.00	0.00	0.00	0.01	0.02	0.04	0.08	0.16	0.30

And if we examine the pledge ratios quantiles, this is also obvious. Failed:

10%	20%	30%	40%	50%	60%	70%	80%	90%
1.01	1.03	1.06	1.10	1.17	1.27	1.46	1.91	3.39

Feature Correlation & Pairs

Kickstarter Data



Correlation of some of the features and overview of their relationship is shown here.

The only items correlated very slightly (.0022) are the **time interval** and the **goal amount in USD**. This makes sense intuitively as the planners expect to raise some amount of money in a specific time.

Functions of Note

This project used several different functions to help with the creation of dynamic model formulas, as well as some wrappers to quickly and neatly get accuracy and confusion matrix data. I'm including these here.

```
results_GLM <- function(train_data, predicted_field_name, predictors_list, model_family= "binomial") {
  fm_string <- paste(predicted_field_name, "~")
  for (i in 1:length(predictors_list)) {
    fm_string <- paste(fm_string, predictors_list[i])
    if (i < length(predictors_list))
      fm_string <- paste(fm_string, "+")}
```

```

}

print(fm_string)
fm <- as.formula(fm_string)

glm_fit <- glm(fm , train_data, family = model_family)
return (glm_fit)
}

get_confusion_matrix <- function(model_fit, test_data)
{
  aug_model <- augment(model_fit, newdata = test_data, type.predict = 'response')

  result <- aug_model %>% mutate(Prediction = factor(round(.fitted)), Reference = state)
    dplyr::select(Reference , Prediction ) %>% table()

  return (caret::confusionMatrix(result))
}

get_accuracy <- function(cf_matrix)
{
  cf_accuracy <- sum(diag(cf_matrix)) / sum(cf_matrix)
  return(cf_accuracy)
}

```

Evaluation Models

For our model calculations and analysis, we will adjust the dataset to remove the variables that are entered at the completion of the project, since using these will skew our calculations. We will only use the variables that a project is created with, as well as some of the features we've engineered from those values.

We will exclude:

- * The final amount of backers for the project.
- * The final amount of money raised for the project. The predictors we use will be chosen from the following fields:
 - * The **Category** the project falls into
 - * The **Currency** the project is raising money in.
 - * The **Goal** raise amount in \$USD
 - * The **Country** the project is based out of
 - * The **Time Interval in Days** the project will last on Kickstarter
 - * The **Day of the Week** the project was launched
 - * The **Month of the Year** the project was launched
 - * The **Year** the project was launched

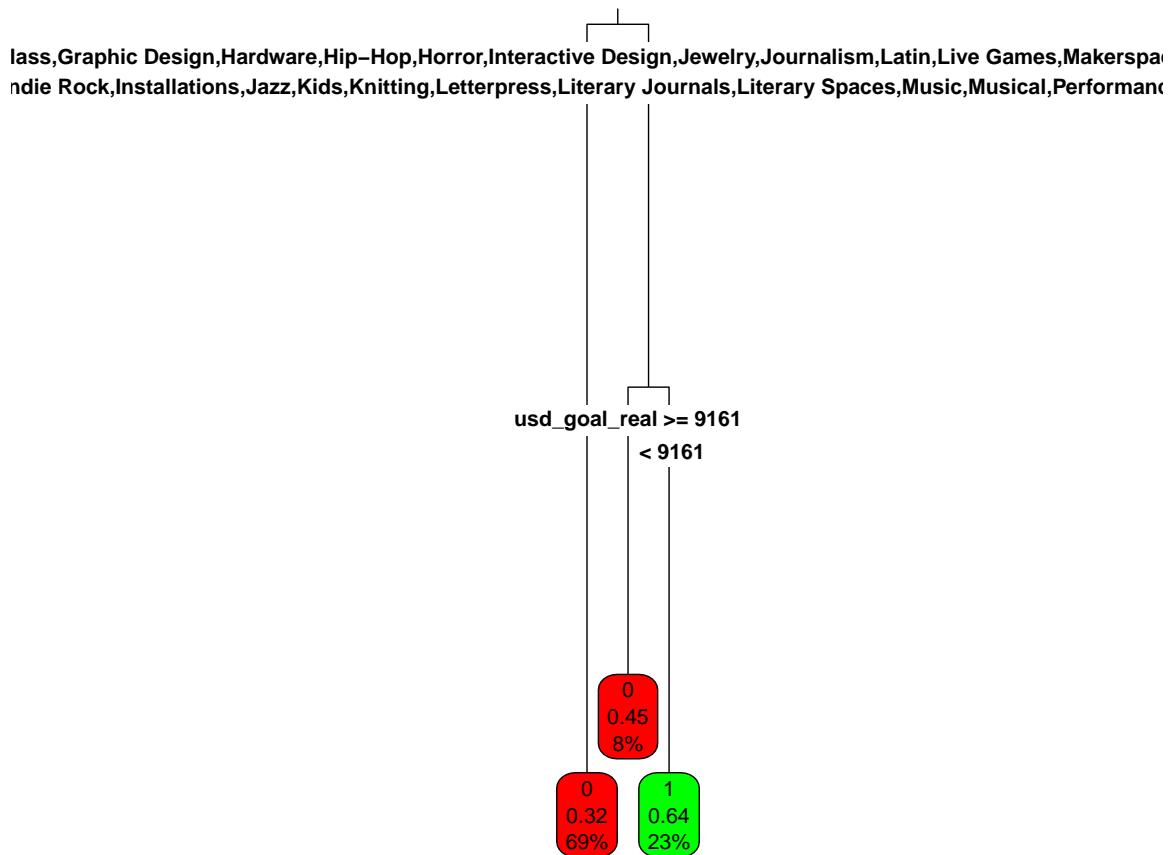
The AI models will try to predict the 'state' of each kickstarter project in the test dataset.

Since we are trying to predict the variable **state** and it is **categorical**, as well as are many of the predictors, **linear regression is not a good option for for this project**, since linear regression only works for numeric variables. Therefor, we attempt the models below.

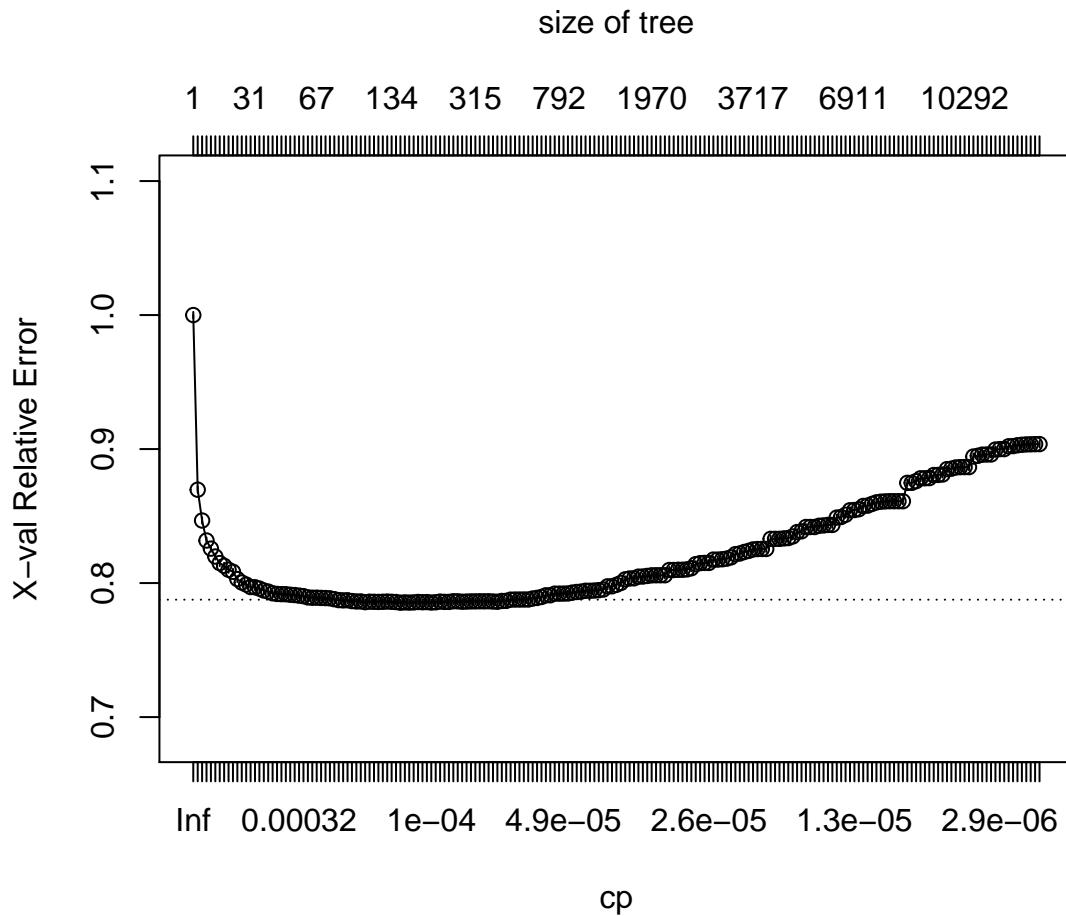
1. Classification Trees

The first model I will attempt will be **classification trees (CT)**. The initial attempt used the following parameters:

1. maxdepth = 5 (the maximum depth of any node of the final tree)
2. minsplit = 200 (the minimum number of observations that must be observed before splitting the tree)



The next attempt once again runs the same model, but this time with **all default parameters**, save **cp** (the complexity parameter), which we set to **0**. The logic behind this is that we make a detailed tree and prune it AFTER the model has run its course. **The plot of the cp values** is shown below, showing the optimal (lowest) value for the relative error.



The code then selects this minimal value, 2.2257502×10^{-4} , and prunes the tree with that complexity parameter. The accuracy of the resulting model is **0.6844813**. As I continue with the project, I will display results of the accuracy evaluations of every model, and print them out as shown below - so that the data can be reviewed as the project progresses.

model	accuracy
Classification Trees mex depth=5,minsplit=200	0.6589737
cp=0 Classification Trees post-pruned with min cp=0.000222575021790562	0.6844813

Next, still attempting to optimize classification trees, I will attempt to run the model with ALL default parameters set, and then again post-prune the resulting tree by finding the minimum **cp** parameter.

This time, the code selects a minimal cp value of **0.01** to prune the tree. The accuracy of the resulting model is **0.6589737**.

model	accuracy
Classification Trees mex depth=5,minsplit=200	0.6589737
cp=0 Classification Trees post-pruned with min cp=0.000222575021790562	0.6844813
Classification Trees, default params, un-pruned	0.6589737
Classification Trees, default params, post-pruned cp=0.01	0.6589737

2. Random Forest

When first attempting to run the **random forest** algorithm on the data, it became obvious that the general implementation wouldn't be plausable here due to some of the factors (the **categories** in particular) having more than 32 levels. (This is a limitation of the random forest algorithm implementation in R.) Therefor, the data had to be converted into a matrix with dummy variables. A sample of 70,000 was taken to evaluate the training set; the processing time for this was very long - and in the future implementation of the model I will attempt to run random forest on a larger subset of data. The final random forest model accuracy with **200 trees and a 70,000 samples** was **0.6866986**.

model	accuracy
Classification Trees mex depth=5,minsplit=200	0.6589737
cp=0 Classification Trees post-pruned with min cp=0.000222575021790562	0.6844813
Classification Trees, default params, un-pruned	0.6589737
Classification Trees, default params, post-pruned cp=0.01	0.6589737
Random Forest, ntree=200	0.6866986

3. Naive Bayes

The next model attempted is the **Naive Bayes Model**. In the initial attempt, I used the same training/testing dataframes as in the classification trees attempts described above. This produced an accuracy of **0.4431321**, which is so far the lowest we've seen (and will prove to be the worst).

model	accuracy
Classification Trees mex depth=5,minsplit=200	0.6589737
cp=0 Classification Trees post-pruned with min cp=0.000222575021790562	0.6844813
Classification Trees, default params, un-pruned	0.6589737
Classification Trees, default params, post-pruned cp=0.01	0.6589737
Random Forest, ntree=200	0.6866986
Naive Bayes	0.4431321

However, the algorithm also takes matrices as input, and at the next iteration the model will attempt to evaluate the accuracy on the matrix data (which is simply the dataset with dummay

variables.) Here, the accuracy was much higher, at **0.5416931**.

model	accuracy
Classification Trees mex depth=5,minsplit=200	0.6589737
cp=0 Classification Trees post-pruned with min cp=0.000222575021790562	0.6844813
Classification Trees, default params, un-pruned	0.6589737
Classification Trees, default params, post-pruned cp=0.01	0.6589737
Random Forest, ntree=200	0.6866986
Naive Bayes	0.4431321
Naive Bayes w/matrix dummy variables	0.5416931

4. Generalized Linear Models (GLM)

The next, and most extensive set of models uses **Generalized Linear Regression** - a variation of linear regression that allows for predictions of categorical variables. Since here we are trying to predict the variable **state**, which is categorical, GLM are a good option to experiment with. I attempt to run the models on several of the built-in and engineered predictors, as shown above. Unsurprisingly, the best-performing model in the GLM category is the one that uses the most factors.

4.1 Predictor - Category

First, we use the sole variable “**category**” as a predictor for the **state** variable. The result is marked as ‘**GLM 1**’, and the calculated accuracy is **0.6507829**.

model	accuracy
Classification Trees mex depth=5,minsplit=200	0.6589737
cp=0 Classification Trees post-pruned with min cp=0.000222575021790562	0.6844813
Classification Trees, default params, un-pruned	0.6589737
Classification Trees, default params, post-pruned cp=0.01	0.6589737
Random Forest, ntree=200	0.6866986
Naive Bayes	0.4431321
Naive Bayes w/matrix dummy variables	0.5416931
GLM 1 predictor	0.6507829

4.2 Predictors - Category and Time In Days

Next, we use the variables “**category**” and **time_int** (the time in days from launch to goal date) as a predictor. This gives us slightly better results (accuracy of **0.6574954**), and is marked as “**GLM 2**”.

model	accuracy
Classification Trees mex depth=5,minsplit=200	0.6589737
cp=0 Classification Trees post-pruned with min cp=0.000222575021790562	0.6844813
Classification Trees, default params, un-pruned	0.6589737
Classification Trees, default params, post-pruned cp=0.01	0.6589737
Random Forest, ntree=200	0.6866986
Naive Bayes	0.4431321
Naive Bayes w/matrix dummy variables	0.5416931
GLM 1 predictor	0.6507829
GLM 2 predictors	0.6574954

4.3 3 Predictors - Category, Time In Days, and Country

Next, we use the variables “**category**” and **time_int**, and **country**. This gives us slightly better results (accuracy of **0.6587172**), and is marked as “**GLM 3**”.

model	accuracy
Classification Trees mex depth=5,minsplit=200	0.6589737
cp=0 Classification Trees post-pruned with min cp=0.000222575021790562	0.6844813
Classification Trees, default params, un-pruned	0.6589737
Classification Trees, default params, post-pruned cp=0.01	0.6589737
Random Forest, ntree=200	0.6866986
Naive Bayes	0.4431321
Naive Bayes w/matrix dummy variables	0.5416931
GLM 1 predictor	0.6507829
GLM 2 predictors	0.6574954
GLM 3 predictors	0.6587172

4.4 4 Predictors - Category, Time In Days, Country and Goal Amount in \$USD

Next, we use the variables “**category**” and **time_int**, **country**, **usd_goal_real**. This gives us slightly better results (accuracy of **0.6678432**), and is marked as “**GLM 4**”.

model	accuracy
Classification Trees mex depth=5,minsplit=200	0.6589737
cp=0 Classification Trees post-pruned with min cp=0.000222575021790562	0.6844813
Classification Trees, default params, un-pruned	0.6589737
Classification Trees, default params, post-pruned cp=0.01	0.6589737
Random Forest, ntree=200	0.6866986
Naive Bayes	0.4431321
Naive Bayes w/matrix dummy variables	0.5416931
GLM 1 predictor	0.6507829
GLM 2 predictors	0.6574954

model	accuracy
GLM 3 predictors	0.6587172
GLM 4 predictors	0.6678432

4.5 5 Predictors - Category, Time In Days, Country, Goal Amount in \$USD, and Launched Month

Next, we use the variables “**category**” and **time_int, country, usd_goal_real**, and “**launched_month**”. This gives us slightly better results (accuracy of **0.6680695**), and is marked as “**GLM 5**”.

model	accuracy
Classification Trees mex depth=5,minsplit=200	0.6589737
cp=0 Classification Trees post-pruned with min cp=0.000222575021790562	0.6844813
Classification Trees, default params, un-pruned	0.6589737
Classification Trees, default params, post-pruned cp=0.01	0.6589737
Random Forest, ntree=200	0.6866986
Naive Bayes	0.4431321
Naive Bayes w/matrix dummy variables	0.5416931
GLM 1 predictor	0.6507829
GLM 2 predictors	0.6574954
GLM 3 predictors	0.6587172
GLM 4 predictors	0.6678432
GLM 5 predictors	0.6680695

4.6 6 Predictors - Category, Time In Days, Country, Goal Amount in \$USD, Launched Month, and Launched Day of the Week

Next, we use the variables “**category**” and **time_int, country, usd_goal_real**, and “**launched_month**”, and **launched_day_of_the_week**. This gives us slightly better results (accuracy of **0.6691556**), and is marked as “**GLM 6**”.

model	accuracy
Classification Trees mex depth=5,minsplit=200	0.6589737
cp=0 Classification Trees post-pruned with min cp=0.000222575021790562	0.6844813
Classification Trees, default params, un-pruned	0.6589737
Classification Trees, default params, post-pruned cp=0.01	0.6589737
Random Forest, ntree=200	0.6866986
Naive Bayes	0.4431321
Naive Bayes w/matrix dummy variables	0.5416931
GLM 1 predictor	0.6507829
GLM 2 predictors	0.6574954
GLM 3 predictors	0.6587172
GLM 4 predictors	0.6678432
GLM 5 predictors	0.6680695

model	accuracy
GLM 6 predictors	0.6691556

4.7 7 Predictors - Category, Time In Days, Country, Goal Amount in \$USD, Launched Month, Launched Day of the Week, and Currency

Next, we use the variables “**category**” and **time_int**, **country**, **usd_goal_real**, “**launched_month**”, and “**currency**”. This gives us the **same** result as ‘GLM 6’ (**0.6691556**) - so it seems that the Currency predictor **wasn’t** helpful. The result is marked as “**GLM 7**”.

model	accuracy
Classification Trees mex depth=5,minsplit=200	0.6589737
cp=0 Classification Trees post-pruned with min cp=0.000222575021790562	0.6844813
Classification Trees, default params, un-pruned	0.6589737
Classification Trees, default params, post-pruned cp=0.01	0.6589737
Random Forest, ntree=200	0.6866986
Naive Bayes	0.4431321
Naive Bayes w/matrix dummy variables	0.5416931
GLM 1 predictor	0.6507829
GLM 2 predictors	0.6574954
GLM 3 predictors	0.6587172
GLM 4 predictors	0.6678432
GLM 5 predictors	0.6680695
GLM 6 predictors	0.6691556
GLM 7 predictors	0.6691556

4.8 8 Predictors - Category, Time In Days, Country, Goal Amount in \$USD, Launched Month, Launched Day of the Week, Currency, and Year Launched.

Next, we use the variables “**category**” and **time_int**, **country**, **usd_goal_real**, “**launched_month**”, “**currency**”, and “**launched_year**”. This gives us slightly better results (**0.6728211**), and is marked as “**GLM 8**”.

model	accuracy
Classification Trees mex depth=5,minsplit=200	0.6589737
cp=0 Classification Trees post-pruned with min cp=0.000222575021790562	0.6844813
Classification Trees, default params, un-pruned	0.6589737
Classification Trees, default params, post-pruned cp=0.01	0.6589737
Random Forest, ntree=200	0.6866986
Naive Bayes	0.4431321
Naive Bayes w/matrix dummy variables	0.5416931
GLM 1 predictor	0.6507829
GLM 2 predictors	0.6574954

model	accuracy
GLM 3 predictors	0.6587172
GLM 4 predictors	0.6678432
GLM 5 predictors	0.6680695
GLM 6 predictors	0.6691556
GLM 7 predictors	0.6691556
GLM 8 predictors	0.6728211

Results

After creating a number of predictive algorithms, the best model as per the project requirements is the one with the highest accuracy so far - this happens to be :

Model - **Random Forest, ntree=200**

Accuracy - **0.6866986**

The models, sorted in descending order of accuracy, are shown below:

model	accuracy
Random Forest, ntree=200	0.6866986
cp=0 Classification Trees post-pruned with min cp=0.000222575021790562	0.6844813
GLM 8 predictors	0.6728211
GLM 6 predictors	0.6691556
GLM 7 predictors	0.6691556
GLM 5 predictors	0.6680695
GLM 4 predictors	0.6678432
Classification Trees mex depth=5,minsplit=200	0.6589737
Classification Trees, default params, un-pruned	0.6589737
Classification Trees, default params, post-pruned cp=0.01	0.6589737
GLM 3 predictors	0.6587172
GLM 2 predictors	0.6574954
GLM 1 predictor	0.6507829
Naive Bayes w/matrix dummy variables	0.5416931
Naive Bayes	0.4431321

Conclusion

Report Summary

In this project, I have built several machine learning algorithm to predict project success prediction system of the Kickstarter dataset, available from Kaggle. Project success, for the purposes of this exercise, was set to be the '**successful**' value of the '**state**' variable of the Kickstarter dataset. For prediction purposes, even though the state variable had 6 different options, we selected projects with only 'succesful' or 'failed' for our prediction engine.

The accuracy of each algorithm has been provided, and tested against a test set, which was selected to be as 20% of the presented data set (the training set being the other 80%). The **Random Forest, ntree=200** is characterized by the **highest accuracy value (0.6866986)** of all the models I have examined, and is hence the optimal model to use for the present project.

Other Options

Exclusion

I have also attemped to run the **K Nearest Neighbors** algorithm on this dataset. While this provided decent results, the processing time to run KNN was un acceptable (> 5 hours.) I have therefore chosen to exclude this model from the project; however, the code is still available (albeit commented out), in the .R file attached to the project.

Possible Extensions and Improvements

The possible evaluation models we can also attempt to get even better results would include:

- Xgboost Model
- GBM Model
- Regularization

Another possiblity for improvement is to be even more creative with feature engineering. Just as we created 3 variables from the launch dates (Month, Day of the Week, and Year launched), we can create the same variables from the deadline variable.

Finally, we can extend our predictions to include not only the categorical **state** variable (that simply predicts the success or failure of the project), and **try to predict the actual amount of moneys raised** for each project.

Appendix - Environment

```
print("Operating System:")  
  
[1] "Operating System:"  
  
version  
  
platform      x86_64-apple-darwin17.0  
arch          x86_64  
os            darwin17.0  
system        x86_64, darwin17.0  
status  
major         4  
minor         0.2  
year          2020  
month         06  
day           22  
svn rev       78730  
language      R  
version.string R version 4.0.2 (2020-06-22)  
nickname      Taking Off Again
```